

Lenov Open Cloud Network Reference Architecture

Feng Xia

Abstract

Lenovo Open Cloud consists of a list of physical servers (aka. nodes) and virtual machines (VMs). This reference environment provides a comprehensive example demonstrating how to set up networks to connect these servers and VMs together.

Contents

1	Introduction	4
2	Business problem and business value	6
2.1	Business problem	6
2.2	Business value	6
3	Architecture Overview	7
4	Hardware	8
4.1	Platform servers	8
4.1.1	Memory	9
4.1.2	RAID controller	9
4.1.3	Disk configurations	10
4.2	Network Switches	10
4.2.1	Lenovo RackSwitch G8272 (10Gb)	10
4.2.2	Lenovo RackSwitch G8052 (1Gb)	10
5	Software Services	12
5.1	Platform services	12
5.1.1	Runtime service	12
5.1.2	Software repository & life cycle management service	12
5.1.3	Automation service	13
5.1.4	Discovery service	13
5.1.5	Inventory planning service	13
5.1.6	Server config & OS deployment service	13
5.1.7	OS image service	13
5.1.8	Configure & Automation repository service	13
5.2	Storage services	13
5.2.1	Ceph capacity management	13
5.3	Cloud services	13
5.3.1	Openstack	13
6	Network Design	14
6.1	Conventions	14
6.2	Connection To Upstream	15
6.3	Platform networks	15
6.3.1	Platform services to VLAN mapping	16
6.3.2	Platform VLAN to server's NIC mapping	17

6.3.3	Platform server's NIC to switch mapping	17
6.4	Storage networks	17
6.4.1	Storage services to VLAN mapping	19
6.4.2	Storage VLAN to server's NIC mapping	19
6.4.3	Storage server's NIC to switch mapping	19
6.5	Cloud networks	20
6.5.1	Cloud VLANs	21
6.5.2	Cloud server's NIC to VLAN mapping	21
6.5.3	Configure Cloud server network interfaces	21
7	Configure Switches	22
7.1	switch to switch	22
7.1.1	cable schema	22
7.1.2	port configurations	22
7.2	switch to server	22
7.2.1	Platform servers	22
7.2.1.1	Platform server to switch cable schema	22
7.2.1.2	Platform server switch port configurations	23
7.2.2	Storage servers	23
7.2.2.1	Storage server to switch cable schema	23
7.2.2.2	Storage server switch port configurations	23
7.2.3	Cloud servers	23
7.2.3.1	Cloud server to switch cable schema	23
7.2.3.2	Cloud server switch port configurations	23
8	Configure Virtual Machines network interfaces	24
9	Appendix	25
9.1	Implementation Worksheet (questionnaire)	25
9.2	Hardware BOM	25
9.3	Software BOM	25
9.3.1	6-server, HCI deployment, 3 year premium	25

List of Figures

3.1	Lenovo Open Cloud Architecture	7
4.1	Lenovo ThinkSystem SR650	8
4.2	ThinkSystem RAID 930-8i RAID Controller	10
4.3	Lenovo RackSwitch G8272	10
4.4	Lenovo RackSwitch G8052	11
6.1	Lenovo Open Cloud Network Overview	14
6.2	Lenovo Open Cloud Platform Network Overview	15
6.3	Lenovo Open Cloud Plaform Service Network Interfaces	16
6.4	Lenovo Open Cloud Plaform Server Network Interfaces	17
6.5	Lenovo Open Cloud Plaform Server to Switch	18
6.6	Lenovo Open Cloud Storage Networks	18
6.7	Lenovo Open Cloud Storage Server Network Interfaces	19
6.8	Lenovo Open Cloud Storage Server to Switch	20

Title: Setting Up Lenovo Open Cloud Networks Date: 2018-09-14 09:35 Tags: reference architecture Slug: ibb network design Author: Feng Xia Status: Draft

Chapter 1

Introduction

The target audience for this Reference Architecture (RA) is system administrators or system architects. Some experience with Red Hat implementation of virtualization, shared storage, and OpenStack is helpful, but it is not required.

Technology evolves fast. We have seen the wave of hardware virtualization in both server space and in personal computing. Then came cloud computing, in which infrastructure becomes even more abstract and remote to end user than ever before. Instead of being viewed as brick and mortar, server, storage, networking are **resources** that can be requested, leased for a period of time, paid per use, and releases when done — all through nothing but an online account and a credit card. The flexibility of this model and the feeling that resource pool can be extended boundlessly has both lowered barrier of entry of new application growing from zero to infinity with little sweat, and elevated requirement on the design, implementation, and operation of such infrastructure

Further more, along the trail of technology evolution, business has been left with an army of legacy systems which were designed and built on a technology stack that was adequate then, but not in trend now. Millions have been invested, millions of users are probably depending on the continuity of service, and many developers and operators were trained and are given the responsibility to maintain such stack. It is neither feasible to cut the cord just because a new technology becomes the talk of the day, nor advisable to continue as before without taking advantage of what new tools can bring. Therefore, it is not only desirable, but in our opinion essential to have an infrastructure that is both flexible and balanced — it must support a broad range of user and application by providing a platform that has a rich mix of building blocks which, first of all, covers common needs out of box, such as keeping an operating system up to date via patch, update and hotfix, while maintaining an open architecture to extend both horizontally in term of resource (compute, storage, networking), and vertically (application stack).

It is with this in mind that Leonov Open Cloud is designed to combine the best of technologies in the market today into a coherent user experience while all the following users will feel at home:

1. **VM users:** Open Cloud supports hardware virtualization in its core. Traditional virtual machines users and applications can be migrated onto the platform while minimizing dependency on underline hardware environment.
2. **Cloud users:** Open Cloud provides on-premise cloud computing environment based on OpenStack, the leading cloud operating system.
3. **Container users:** Devops have continuously pushed the boundary to merge development and

production into a single, consistent experience that what developer uses as a **sandbox** should be identical as what can be used in production. By doing so not only we will eliminate the necessity to maintain multiple stacks catering for different environments — a typical setup will be one for development, one for testing, and one for production, but minimize chance of incompatibility and bugs due to difference between two environments.

This RA describes the system architecture for the Lenovo Open Cloud Platform based on Lenovo ThinkSystem servers and Lenovo network switches. It provides detail of the hardware requirements to support various node roles and the corresponding configuration of the systems. It also describes the network architecture and details for the switch configurations. The hardware bill of materials is provided for all required components to build the Open Cloud cluster. An example deployment is used to show how to prepare, provision, deploy, and manage the Open Cloud on Lenovo ThinkSystem servers and Lenovo network switches.

Chapter 2

Business problem and business value

2.1 Business problem

2.2 Business value

Chapter 3

Architecture Overview

Lenovo Open Cloud has two sets of clusters: **management cluster** and **workload cluster**. Workload cluster refers to applications directly interfacing with end user. Management cluster refers to applications that manage and provide Open Cloud services.

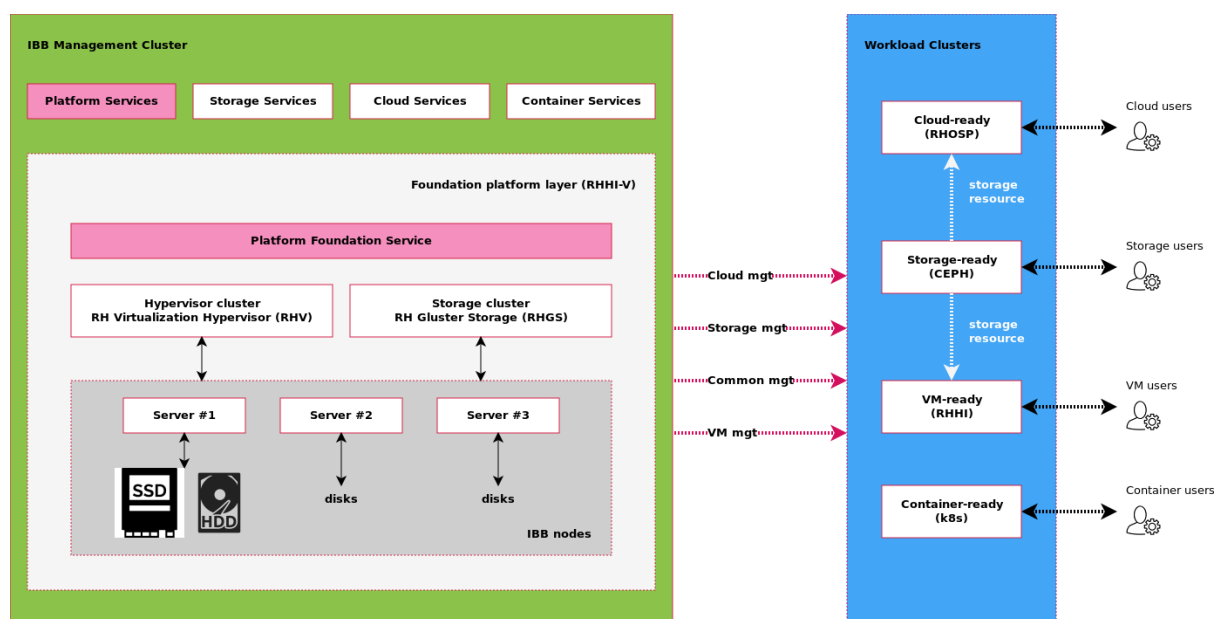


Figure 3.1: Lenovo Open Cloud Architecture

Management cluster includes a **platform foundation layer** and four groups of service.

Platform foundation Foundation layer includes 3-, 6-, or 9- servers depending on configuration, storage disks inside server, two 1Gb switches, and two 10Gb switches. It supports virtual machines on top of Red Hat Virtualization. Storage uses Gluster FS across servers.

Chapter 4

Hardware

Lenovo Open Cloud is a highly configurable system. From the point of view of physical servers, LOC can be deployed on a 6-server or 9-server configuration. In this document we will use a 9-server configuration as example.

4.1 Platform servers



Figure 4.1: Lenovo ThinkSystem SR650

The Lenovo ThinkSystem SR650 server is an enterprise class 2U two-socket versatile server that incorporates outstanding reliability, availability, and serviceability (RAS), security, and high efficiency for business-critical applications and cloud deployments. Unique Lenovo AnyBay technology provides the flexibility to mix-and-match SAS/SATA HDDs/SSDs and NVMe SSDs in the same drive bays. Four direct-connect NVMe ports on the motherboard provide ultra-fast read/writes with NVMe drives and

reduce costs by eliminating PCIe switch adapters. Plus, storage can be tiered for greater application performance, to provide the most cost-effective solution.

Combined with the Intel® Xeon® Scalable processors product family, the Lenovo ThinkSystem SR650 server offers a high density combination of workloads and performance. Its flexible, pay-as-you-grow design and great expansion capabilities solidify dependability for any kind of virtualized workload, with minimal downtime. Additionally, it supports two 300W high-performance GPUs and ML2 NIC adapters with shared management.

The Lenovo ThinkSystem SR650 server provides internal storage density of up to 100 TB (with up to 26 x 2.5-inch drives) in a 2U form factor with its impressive array of workload-optimized storage configurations. The ThinkSystem SR650 offers easy management and saves floor space and power consumption for the most demanding storage virtualization use cases by consolidating the storage and server into one system. The Lenovo ThinkSystem SR650 server supports up to twenty-four 2.5-inch or fourteen 3.5-inch hot-swappable SAS/SATA HDDs or SSDs together with up to eight on-board NVMe PCIe ports that allow direct connections to the U.2 NVMe PCIe SSDs. The ThinkSystem SR650 server also supports up to two NVIDIA GRID cards for AI or media processing acceleration.

The SR650 server supports up to two processors, each with up to 28-core or 56 threads with hyper-threading enabled, up to 38.5 MB of last level cache (LLC), up to 2666 MHz memory speeds and up to 3 TB of memory capacity. The SR650 also support up to 6 x PCIe slots. Its on-board Ethernet solution provides 2/4 standard embedded Gigabit Ethernet ports and 2/4 optional embedded 10 Gigabit Ethernet ports without occupying PCIe slots. All these advanced features make the server ideal to run data and bandwidth intensive VNF workload and storage functions of NFVI platform.

For more information, see [product guide](#).

4.1.1 Memory

There are 24 slots in total for memory in the SR650 server. The maximum memory for each slot is 128 GB. So, the maximum memory of one SR650 server can reach $128\text{GB} * 24 \text{ slots} = 3\text{TB}$.

For each SR650 server, it is recommended to use 64GB memory for a small deployment, 128G memory for a medium deployment, and 256GB for a large deployment.¹

4.1.2 RAID controller

The ThinkSystem RAID 930 family of internal 12 Gbps SAS RAID controllers are high-performance RAID-on-chip (ROC) adapters. These adapters support RAID levels 0/1/10/5/50/6/60 as well as JBOD, and include an extensive list of RAS and management features.

The family is comprised of four adapters:

1. The ThinkSystem RAID 930-4i supports up to four internal SAS and SATA drives
2. The ThinkSystem RAID 930-8i supports up to eight internal SAS and SATA drives
3. The ThinkSystem RAID 930-16i supports up to 16 internal SAS and SATA drives
4. The ThinkSystem RAID 930-24i supports up to 24 internal SAS and SATA drives

For more information, see [product guide](#).

¹See Red Hat [DEPLOYING RED HAT HYPERCONVERGED INFRASTRUCTURE](#) for details.




Figure 4.2: ThinkSystem RAID 930-8i RAID Controller

4.1.3 Disk configurations

Disk configuration is important to achieve high performance. We recommend the following in each SR650 server:

Type	Position	Number	Size	RAID	Purpose
SSD	Front backplane	2	800GB	RAID1	Operating system
SSD	Front backplane	2	128GB	RAID1	LVM cache
SAS HDD	Rear backplane	8	2TB	RAID6	Glusterfs

Total Size for Gluster: $3 * 8 * 2TB = 48TB$.

4.2 Network Switches

The following sections describe the Top-of-Rack (ToR) switches used in this reference architecture. The Networking Operating System software features of these Lenovo switches deliver seamless, standards-based integration into upstream switches. Two 10 Gb switches and two 1Gb switches are used in this architecture.

4.2.1 Lenovo RackSwitch G8272 (10Gb)



Figure 4.3: Lenovo RackSwitch G8272

The Lenovo RackSwitch G8272 uses 10Gb SFP+ and 40Gb QSFP+ Ethernet technology and is specifically designed for the data center. It is an enterprise class Layer 2 and Layer 3 full featured switch that delivers line-rate, high-bandwidth, low latency switching, filtering, and traffic queuing without delaying data. Large data center-grade buffers help keep traffic moving, while the hot-swap redundant power supplies and fans (along with numerous high-availability features) help provide high availability for business sensitive traffic.



Figure 4.4: Lenovo RackSwitch G8052

RJ-45 ports and four 10 Gigabit Ethernet SFP+ ports (it also supports 1 GbE SFP transceivers), and includes hot-swap redundant power supplies and fans as standard, which minimizes your configuration requirements. Unlike most rack equipment that cools from side-to-side, the G8052 has rear-to-front or front-to-rear airflow that matches server airflow.

For more information, see [product guide](#).

Chapter 5

Software Services

Lenovo Open Cloud software can be viewed in three categories based on the services they provide:

1. **Platform services:** Platform services are built upon [Red Hat Hyperconverged Infrastructure \(RHHI\)](#). It provides LOC core services each deployed in one or more virtual machines.
2. **Storage services:** Storage services are built upon Ceph. It provides capability to manage Ceph cluster up to xx.
3. **Cloud services:** Cloud is built upon Red Hat Openstack.

5.1 Platform services

Platform services provide administrative functions to support operation of the Open Cloud. This includes management of software life cycle, automation, list of artifacts such as ISO images and qcow images, and new server discovery.

5.1.1 Runtime service

Built upon [Red Hat Hyperconverged Infrastructure \(RHVI\)](#). It supports virtual machine users out of box, and is the foundation of other Lenovo Open Cloud services.

RHHI integrates Red Hat Virtualization (RHV) and Red Hat Gluster Storage (RHGS). RHHI for Virtualization provides open-source, centrally administered, and cost-effective integrated compute and storage in a compact footprint for remote sites.

See [product guide](#) for details.

5.1.2 Software repository & life cycle management service

Built upon [Red Hat Satellite](#). All Open Cloud servers are registered to this service, who then is responsible to manage life cycle of:

1. RHEL and Red Hat software products that are deployed in the Open Cloud.
2. Release, update, patch of Lenovo software products.
3. `.iso` and `qcow2` images, which are used by VM creation and server provisioning.

Satellite is an on-premise alternative to trying to download all of your content from the Red Hat content delivery network or managing your subscriptions through the Customer Portal. From a performance side, it reduces hits to your network bandwidth because local systems can download everything they need locally; from a security side, it can limit the risks of malicious content or access, even enabling entirely disconnected environments.

Satellite is composed of a centralized Satellite Server. Depending on your data center setup, organization design, and geographic locations, you can have local Capsule Servers, which are proxies that locally manage content and obtain subscription, registration, and content from the central Satellite Server.

See [product guide](#) for details.

5.1.3 Automation service

Build upon [Red Hat Ansible Tower](#). It is the single point of contact to manage servers and VMs using ansible playbooks.

Lenovo Open Cloud is shipped with a list of pre-defined automations that makes managing the infrastructure easy and efficient.

See [product guide](#) for details.

5.1.4 Discovery service

Build upon [Lenovo Confluent](#). It continuously monitors network for new Lenovo server and switch. Once identified, the new hardware can be enlisted by other Open Cloud services, such as extending Ceph cluster or adding an Openstack compute node.

See [product guide](#) for details.

5.1.5 Inventory planning service

5.1.6 Server config & OS deployment service

5.1.7 OS image service

5.1.8 Configure & Automation repository service

5.2 Storage services

5.2.1 Ceph capacity management

5.3 Cloud services

5.3.1 Openstack

Chapter 6

Network Design

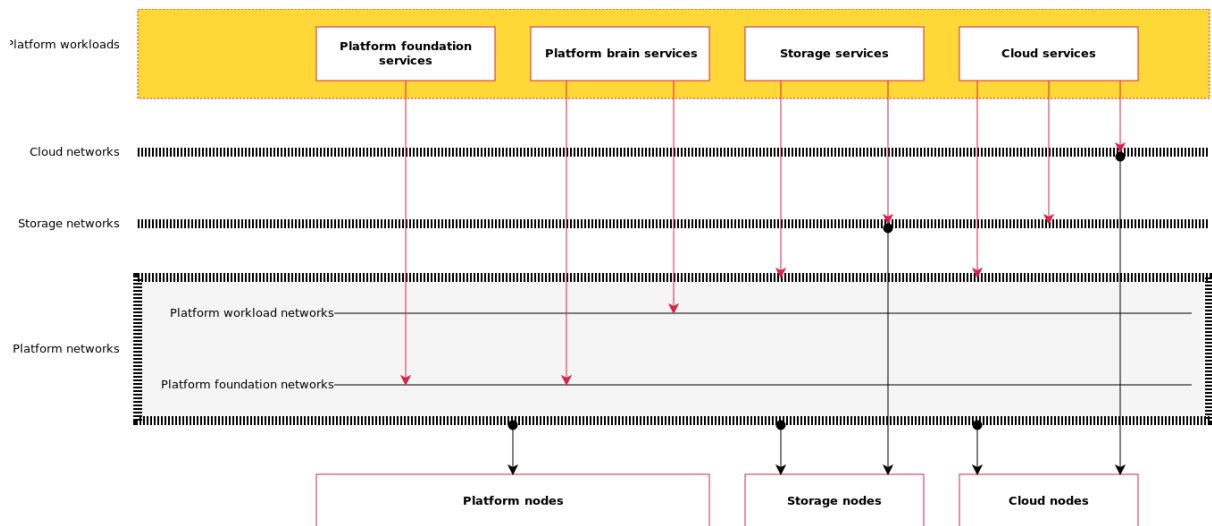


Figure 6.1: Lenovo Open Cloud Network Overview

LOC networks can be viewed in three groups whereas:

1. **platform network:** to support platform services.
2. **storage network:** built on top of platform network with added networks to handle Ceph data storage traffic and Ceph management functions.
3. **cloud network:**

6.1 Conventions

Hardware can break. It is important to keep this in mind when designing a network connection. In this architecture we have followed these conventions:

1. Inter switch connections are paired.
2. Except BMC connection, server to switch connections are paired.
 1. Each pair connect to separate NICs on the server at north bound, and separate switch at south bound.

This then requires matching configuration on the switch using LACP, and on the server using **active-active bonding**.

6.2 Connection To Upstream

Showing switch topology within IBB as well as how it is connected to upstream → what is required from upstream, eg. dhcp, dns, gateway, access to RH CDN.

6.3 Platform networks

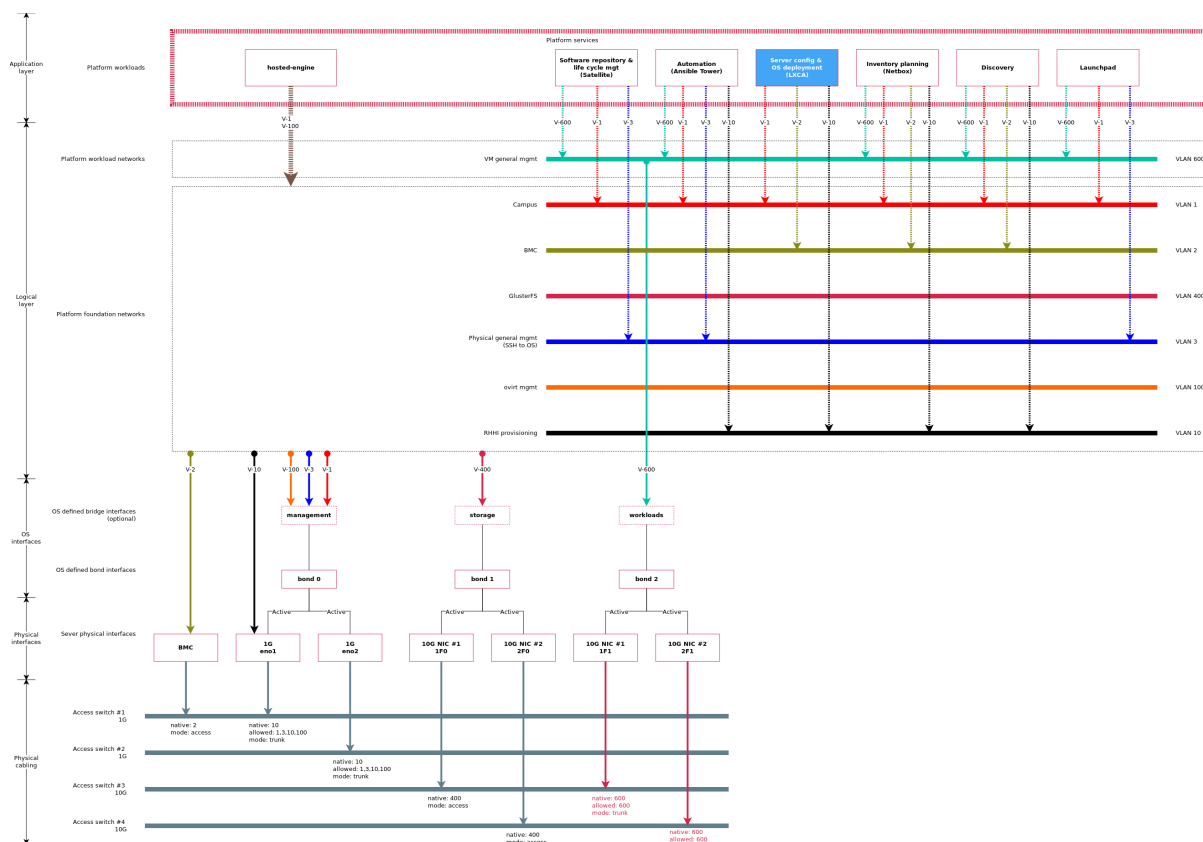


Figure 6.2: Lenovo Open Cloud Platform Network Overview

Platform networks are designed to offer performance and high availability. For illustration purpose we are to separate networks by their function so to highlight some design considerations. It is possible to merge these to fewer networks, or to reuse existing ones for the purpose. For this use section “VLAN Mapping Worksheet”.

Campus campus network is a name for public access. This is the network that Open Cloud user uses to access its services, eg. Ansible Tower’s web UI.

BMC also called **Out-of-Band (OOB)** management network. It connects to a dedicated management port on physical server that is separated from data ports.

ovirt management is a private network linking RHHI management console to RHHI clusters. Except Platform admin, other users should not have direct access to it.

glusterFS is a private network used by **gluster** clusters. Gluster cluster is the storage backend of Open Cloud Platform. For example, in a 6-server configuration, three platform servers will form a 3-node gluster cluster.

Physical server management is to support traffic of In-Band managerial tasks, eg. **ssh** to a server.

OS provisioning is to support data traffic of installing OS on a physical server. Separating this to its own network is a best practice because operating system image can be large, thus its loading to server can have negative impact on shared traffics.

VM management is to access RHHI virtual machines. This supports both the Open Cloud services and VM workloads. Later we will see that it's also advised to dedicate a NIC for this same purpose.

Network	VLAN	Subnet	Addresses	Mask	Static / DHCP	Gateway
Campus	1	10.240.x.x ¹	10		static	10.240.x.1
BMC	2	192.168.2.x	254	/24	static	192.168.2.1
Physical server management	3	192.168.3.x	254	/24	static	192.168.3.1
OS provisioning	10	192.168.10.x	3/6/9	/24	static	192.168.10.1
OVRT management	100	192.168.100.x	3/6/9	/29	static	192.168.100.1
glusterFS	400	192.168.40.x	3/6/9	/29	static	192.168.40.1
VM management	600	192.168.60.x	11	/28	static	192.168.60.1

6.3.1 Platform services to VLAN mapping

Platform Services	1	2	3	10	100	400	600
Runtime	x				x		
Software repository & life cycle management	x		x				x
Automation	x		x	x			x
Server config & OS deployment	x	x		x			
Inventory planning	x	x		x			x
Discovery	x	x		x			x
OS image							
Configure & Automation repository service							




Figure 6.3: Lenovo Open Cloud Platform Service Network Interfaces

¹This is an example subnet.

6.3.2 Platform VLAN to server's NIC mapping

Besides BMC port, each server has minimal two 1Gb ports and four 10Gb ports. Interfaces are paired to form an active-active bonding interface on Platform server. Optionally, we can also create a network bridge on top of a bonding interface. A sample configuration is shown below:

Network	VLAN	BMC	2 x 1G	2 x 10G	2 x 10G	Bond	Bridge
Campus	1		x			0	management
BMC	2	x				n/a	n/a
Physical server management	3		x			0	management
OS provisioning	10		x ²			n/a	
OVRT management	100		x			0	management
glusterFS	400			x		1	storage
VM management	600				x	2	workloads

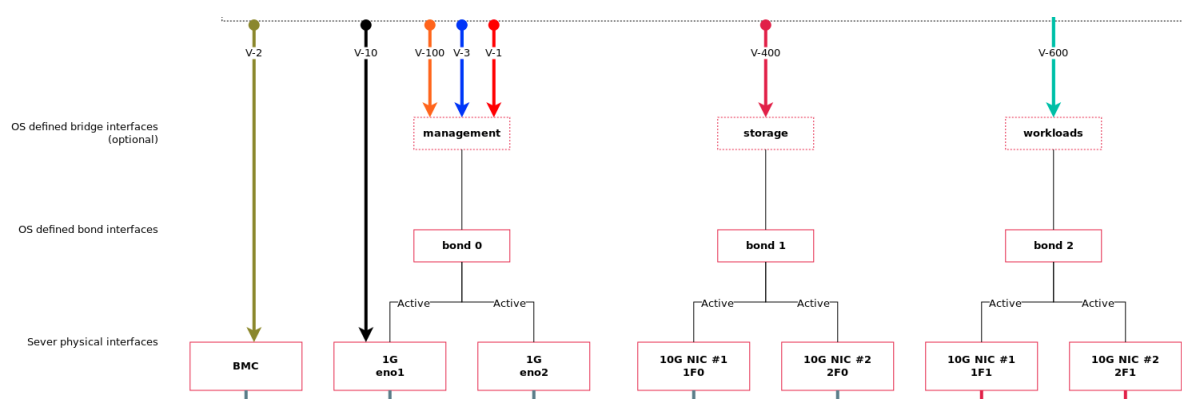


Figure 6.4: Lenovo Open Cloud Platform Server Network Interfaces

6.3.3 Platform server's NIC to switch mapping

Server Side	Switch port mode	Switch port native VLAN	Switch port tagged VLAN
BMC	access	2	n/a
2 x 1G	trunk	10	1,3,10,100
2 x 10G	access	400	n/a
2 x 10G	trunk	600	600

6.4 Storage networks

Lenovo Open Cloud supports Ceph storage backend. A storage backend can be shared among multiple workloads and platforms, such as OpenStack. Three new networks are added to [Platform Network](#) for Ceph function while leveraging [platform services](#) to support Storage hardware and software workloads.

Ceph management is communication between Ceph dashboard and Ceph nodes, e.g. RPC, transferring

²Bonding provisioning network is optional because loading an operating system is not a frequent event.

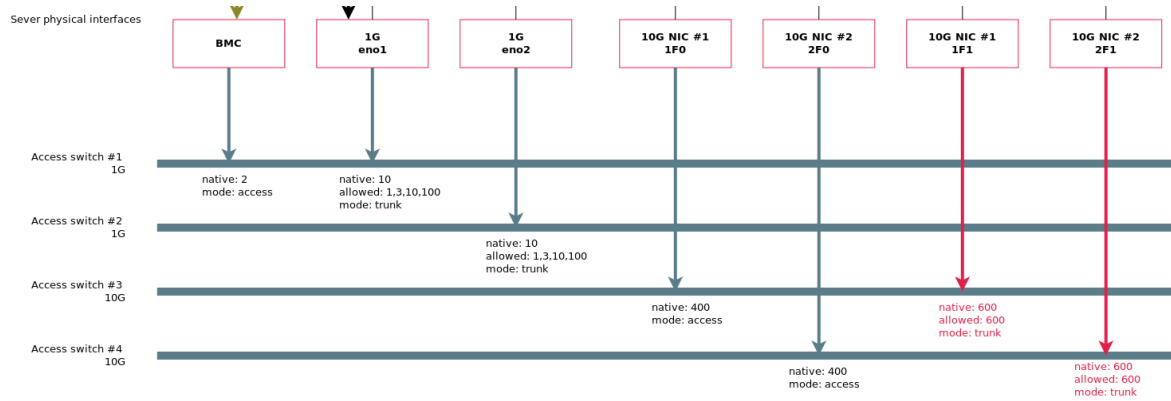


Figure 6.5: Lenovo Open Cloud Platform Server to Switch

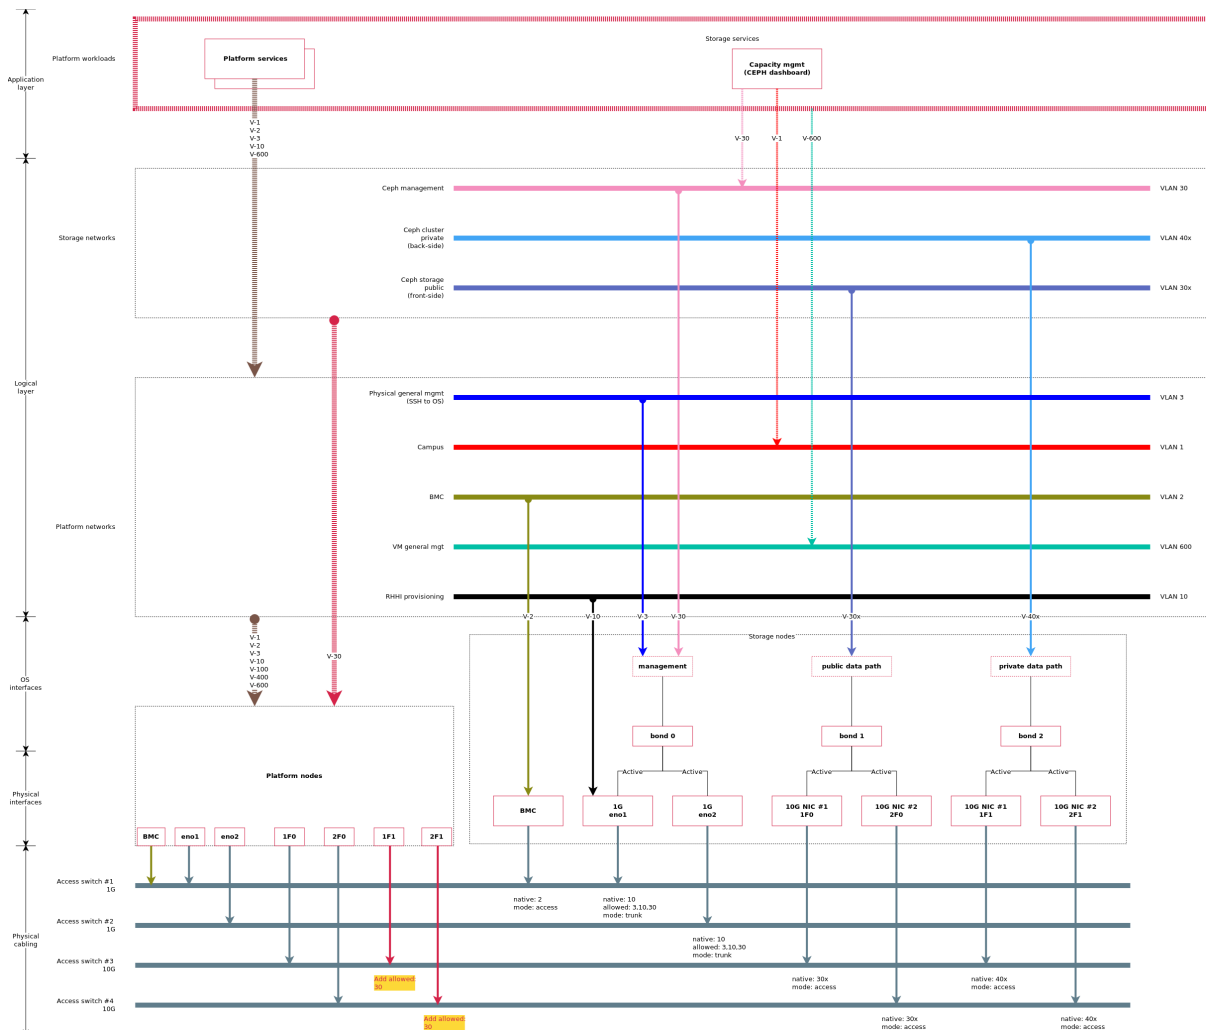


Figure 6.6: Lenovo Open Cloud Storage Networks

zabbix monitoring data.

Ceph storage public Data transferring between OpenStack nodes (both controller nodes and compute nodes) and Ceph nodes.

Ceph cluster private Ceph private data transferring, e.g. rebalancing.

Network	VLAN	Subnet	Addresses	Mask	Static / DHCP	Gateway
Ceph management	30	192.168.30.x	254	/24	static	192.168.30.1
Ceph storage public	30x	192.168.3.x	254	/24	static	192.168.3.1
Ceph cluster private	40x	192.168.4.x	254	/24	static	192.168.4.1

6.4.1 Storage services to VLAN mapping

Storage Services	1	30	600
Capacity management	x	x	x

6.4.2 Storage VLAN to server's NIC mapping

Network	VLAN	BMC	2 x 1G	2 x 10G	2 x 10G	Bond	Bridge
BMC	2	x				n/a	n/a
Ceph management	30		x			0	management
Ceph storage public	30x			x			public data path
Ceph cluster private	40x				x		private data path

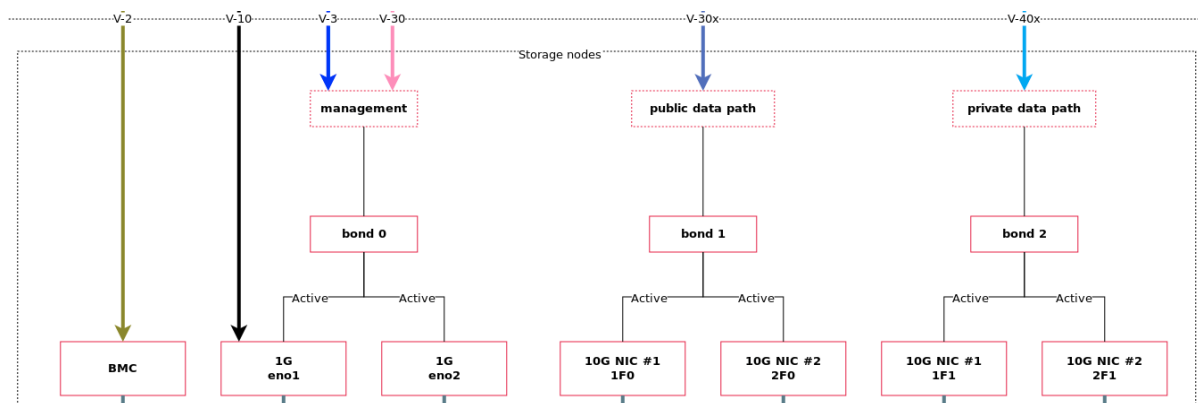


Figure 6.7: Lenovo Open Cloud Storage Server Network Interfaces

6.4.3 Storage server's NIC to switch mapping

Server Side	Switch port mode	Switch port native VLAN	Switch port tagged VLAN
BMC	access	2	n/a
2 x 1G	trunk	10	3,10,30

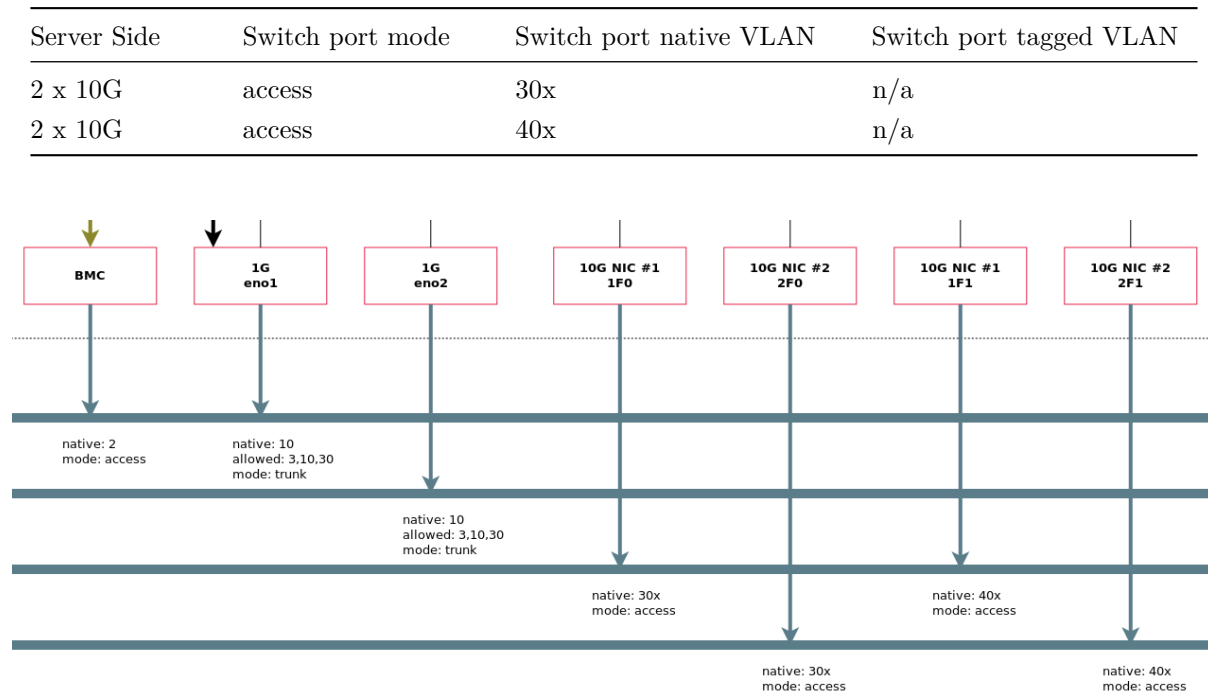
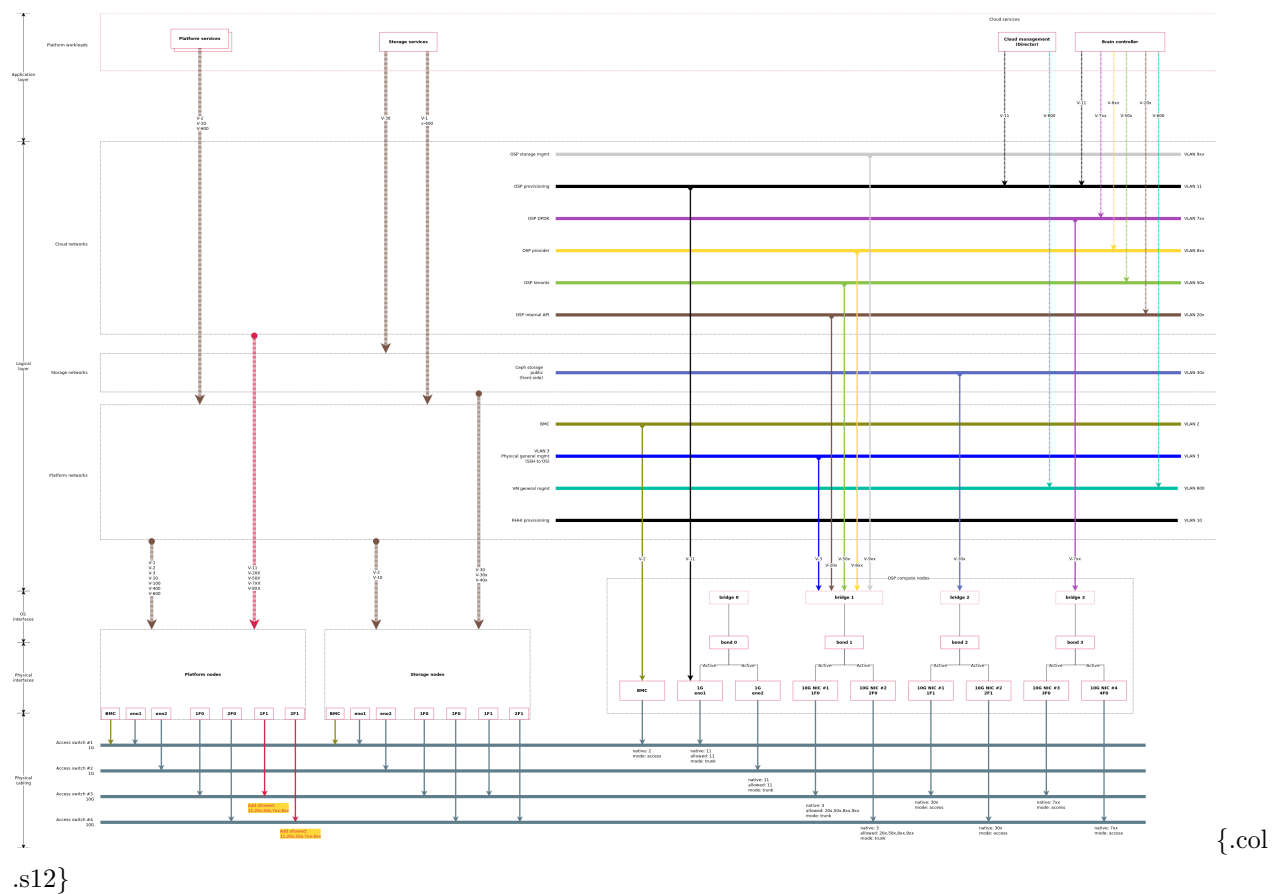


Figure 6.8: Lenovo Open Cloud Storage Server to Switch

6.5 Cloud networks



6.5.1 Cloud VLANs**6.5.2 Cloud server's NIC to VLAN mapping****6.5.3 Configure Cloud server network interfaces**

Chapter 7

Configure Switches

There are two aspects of switch configurations:

1. **inter switch connections:** switch are connected to form a topology allowing data traffic between Lenovo Open Cloud environment and its host environments, and between LOC switches within the LOC itself. All switches are paired for high availability.
2. **server connections:** are connections between server and switch. Except out-of-band connection which has only one connection between a server and a switch, thus does not have redundancy, all other server to switch connections are in pairs.

7.1 switch to switch

7.1.1 cable schema

7.1.2 port configurations

7.2 switch to server

7.2.1 Platform servers

7.2.1.1 Platform server to switch cable schema

Each environment is different. Here we present an example cable schema following the network designs laid out in previous sections. In the following sections we will use this schema to demonstrate switch port configurations.

Port	1G switch	1G switch	10G switch	10G switch
1	server 1 BMC		server 1 bond 1	server 1 bond 1
2	server 2 BMC		server 2 bond 1	server 2 bond 1
3	server 3 BMC		server 3 bond 1	server 3 bond 1
17	server 1 bond 0	server 1 bond 0		
18	server 2 bond 0	server 2 bond 0		
19	server 3 bond 0	server 3 bond 0		

Port	1G switch	1G switch	10G switch	10G switch
------	-----------	-----------	------------	------------

7.2.1.2 Platform server switch port configurations

7.2.2 Storage servers

7.2.2.1 Storage server to switch cable schema

7.2.2.2 Storage server switch port configurations

7.2.3 Cloud servers

7.2.3.1 Cloud server to switch cable schema

7.2.3.2 Cloud server switch port configurations

Chapter 8

Configure Virtual Machines network interfaces

Chapter 9

Appendix

9.1 Implementation Worksheet (questionnaire)

9.2 Hardware BOM

Simplified version covers server & switch at high level should be fine.

9.3 Software BOM

9.3.1 6-server, HCI deployment, 3 year premium

Table 9.1: Software BOM, 6-server, HCI deployment, 3 year premium

SKU	Product	Qty
RS00139F3	Red Hat Hyperconverged Infrastructure for Virtualization (RHHI-V)	1
MCT3305F3	Red Hat Ansible Tower	1
MCT2981F3	Red Hat Openstack Platform (w/o Guest) with Smart Management	4
MCT2979F3	Red Hat OpenStack Platform with Smart Management & Guests	3
RS00036F3	Red Hat Ceph Storage	1
RS00031F3	Smart Management	3
MCT2838F3	Cloudforms	1
MCT3474F3	Red Hat Insights	1